

Feature_Engineering_Lab

May 5, 2022

1 Feature Engineering

Estimated time needed: **45** minutes

A critical part of the successful Machine Learning project is coming up with a good set of features to train on. This process is called feature engineering, and it involves three steps: feature transformation (transforming the original features), feature selection (selecting the most useful features to train on), and feature extraction (combining existing features to produce more useful ones). In this notebook we will explore different tools in Feature Engineering.

1.1 Objectives

After completing this lab you will be able to:

- Understand the types of Feature Engineering
 - Feature Transformation
 - * Dealing with Categorical Variables
 - One Hot Encoding
 - Label Encoding
 - * Date Time Transformations
 - Feature Selection
 - Feature Extraction using Principal Component Analysis
-

1.2 Setup

For this lab, we will be using the following libraries:

- [pandas](#) for managing the data.
- [numpy](#) for mathematical operations.
- [seaborn](#) for visualizing the data.
- [matplotlib](#) for visualizing the data.
- [plotly.express](#) for visualizing the data.
- [sklearn](#) for machine learning and machine-learning-pipeline related functions.

1.3 Installing Required Libraries

The following required modules are pre-installed in the Skills Network Labs environment. However if you run this notebook commands in a different Jupyter environment (e.g. Watson Studio or

Ananconda) you will need to install these libraries by removing the # sign before !mamba in the code cell below.

```
[ ]: # All Libraries required for this lab are listed below. The libraries
      ↪ pre-installed on Skills Network Labs are commented.
      # !mamba install -qy pandas==1.3.4 numpy==1.21.4 seaborn==0.9.0 matplotlib==3.5.
      ↪ 0 scikit-learn==0.20.1
      # Note: If your environment doesn't support "!mamba install", use "!pip install"
```

```
[1]: !mamba install -qy openpyxl
```

Package	Version	Build	Channel	Size
Install:				
+ et_xmlfile	1.1.0	py37h06a4308_0	pkgs/main/linux-64	10kB
+ openpyxl	3.0.9	pyhd3eb1b0_0	pkgs/main/noarch	168kB
Upgrade:				
- ca-certificates	2021.10.8	ha878542_0	installed	
+ ca-certificates	2022.4.26	h06a4308_0	pkgs/main/linux-64	127kB
Summary:				
Install: 2 packages				
Upgrade: 1 packages				
Total download: 304kB				

```
Preparing transaction: ...working... done
Verifying transaction: ...working... done
Executing transaction: ...working... done
```

```
[2]: # Surpress warnings from using older version of sklearn:
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
```

```
[3]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import plotly.express as px

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
```

1.4 Reading and understanding our data

For this lab, we will be using the `airlines_data.xlsx` file, hosted on IBM Cloud object. This dataset contains the prices of flight tickets for various airlines between the months of March and June of 2019 and between various cities. This dataset is often used for prediction analysis of the flight prices which are influenced by various factors, such as name of the airline, date of journey, route, departure and arrival times, the source and the destination of the trip, duration and other parameters.

In this notebook, we will use the airlines dataset to perform feature engineering on some of its independent variables.

Let's start by reading the data into *pandas* data frame and looking at the first 5 rows using the `head()` method.

```
[4]: data = pd.read_excel('https://cf-courses-data.s3.us.cloud-object-storage.
↳ appdomain.cloud/IBM-ML0232EN-SkillsNetwork/asset/airlines_data.xlsx')
data.head()
```

```
[4]:
```

	Airline	Date_of_Journey	Source	Destination	Route	\
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL	
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR	
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL → LKO → BOM → COK	
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU → NAG → BLR	
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR → NAG → DEL	

	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	Price
0	22:20	01:10 22 Mar	2h 50m	non-stop	No info	3897
1	05:50	13:15	7h 25m	2 stops	No info	7662
2	09:25	04:25 10 Jun	19h	2 stops	No info	13882
3	18:05	23:30	5h 25m	1 stop	No info	6218
4	16:50	21:35	4h 45m	1 stop	No info	13302

By using the `info` function, we will take a look at the types of data that our dataset contains.

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
```

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Airline	10683 non-null	object
1	Date_of_Journey	10683 non-null	object
2	Source	10683 non-null	object
3	Destination	10683 non-null	object
4	Route	10682 non-null	object
5	Dep_Time	10683 non-null	object
6	Arrival_Time	10683 non-null	object
7	Duration	10683 non-null	object
8	Total_Stops	10682 non-null	object
9	Additional_Info	10683 non-null	object
10	Price	10683 non-null	int64

dtypes: int64(1), object(10)

memory usage: 918.2+ KB

As we see from the output above, we mostly have object data types, except for the 'price' column, which is an integer.

The `describe()` function provides the statistical information about the numerical variables. In our case, it is the 'price' variable.

```
[6]: data.describe()
```

```
[6]:          Price
count  10683.000000
mean    9087.064121
std     4611.359167
min     1759.000000
25%     5277.000000
50%     8372.000000
75%    12373.000000
max     79512.000000
```

Next, we will check for any null values.

```
[7]: data.isnull().sum()
```

```
[7]: Airline          0
     Date_of_Journey  0
     Source          0
     Destination     0
     Route           1
     Dep_Time        0
     Arrival_Time    0
     Duration        0
     Total_Stops     1
     Additional_Info  0
```

```
Price          0
dtype: int64
```

Now that we have found some null points, we need to either remove them from our dataset or fill them with something else. In this case, we will use `fillna()` and `method='ffill'`, which fills the last observed non-null value forward until another non-null value is encountered.

```
[8]: data = data.fillna(method='ffill')
```

1.5 Feature Transformation

Feature Transformation means transforming our features to the functions of the original features. For example, feature encoding, scaling, and discretization (the process of transforming continuous variables into discrete form, by creating bins or intervals) are the most common forms of data transformation.

1.5.1 Dealing with Categorical Variables

Categorical variables represent qualitative data with no apparent inherent mathematical meaning. Therefore, for any machine learning analysis, all the categorical data must be transformed into the numerical data types. First, we'll start with 'Airlines' column, as it contains categorical values. We will use `unique()` method to obtain all the categories in this column.

```
[9]: data['Airline'].unique().tolist()
```

```
[9]: ['IndiGo',
      'Air India',
      'Jet Airways',
      'SpiceJet',
      'Multiple carriers',
      'GoAir',
      'Vistara',
      'Air Asia',
      'Vistara Premium economy',
      'Jet Airways Business',
      'Multiple carriers Premium economy',
      'Trujet']
```

From the above list, we notice that some of the airline names are being repeated. For example, 'Jet Airways' and 'Jet Airways Business'. This means that some of the airlines are subdivided into separate parts. We will combine these 'two-parts' airlines to make our categorical features more consistent with the rest of the variables.

Here, we will use the *numpy* `where()` function to locate and combine the two categories.

```
[10]: data['Airline'] = np.where(data['Airline']=='Vistara Premium economy',
    ↪ 'Vistara', data['Airline'])
data['Airline'] = np.where(data['Airline']=='Jet Airways Business', 'Jet',
    ↪ 'Airways', data['Airline'])
```

1.6 Exercise 1

In this exercise, use `np.where()` function to combine ‘Multiple carriers Premium economy’ and ‘Multiple carriers’ categories, like shown in the code above. Print the newly created list using `unique().tolist()` functions.

```
[12]: # Enter your code and run the cell
data['Airline'] = np.where(data['Airline'] == 'Multiple carriers Premium_
    ↪economy', 'Multiple carriers', data['Airline'])
data['Airline'].unique().tolist()
```

```
[12]: ['IndiGo',
      'Air India',
      'Jet Airways',
      'SpiceJet',
      'Multiple carriers',
      'GoAir',
      'Vistara',
      'Air Asia',
      'Trujet']
```

Solution (Click Here)

    <code>

```
data['Airline'] = np.where(data['Airline']=='Multiple carriers Premium economy', 'Multiple carriers', data['Airline']) data['Airline'].unique().tolist()
```

One Hot Encoding Now, to be recognized by a machine learning algorithms, our categorical variables should be converted into numerical ones. One way to do this is through *one hot encoding*. To learn more about this process, please visit this [documentation](#).

We will use, `get_dummies()` method to do this transformation. In the next cell, we will transform ‘Airline’, ‘Source’, and ‘Destination’ into their respective numeric variables. We will put all the transformed data into a ‘data1’ data frame.

```
[13]: data1 = pd.get_dummies(data=data, columns = ['Airline', 'Source',
    ↪'Destination'])
```

```
[14]: data1.head()
```

```
[14]:  Date_of_Journey      Route Dep_Time  Arrival_Time  Duration  \
0      24/03/2019      BLR → DEL    22:20    01:10 22 Mar    2h 50m
1      1/05/2019  CCU → IXR → BBI → BLR    05:50         13:15    7h 25m
2      9/06/2019  DEL → LKO → BOM → COK    09:25    04:25 10 Jun      19h
3      12/05/2019      CCU → NAG → BLR    18:05         23:30    5h 25m
4      01/03/2019      BLR → NAG → DEL    16:50         21:35    4h 45m

   Total_Stops  Additional_Info  Price  Airline_Air Asia  Airline_Air India  \
0      non-stop           No info   3897              0              0
```

1	2 stops	No info	7662	0	1
2	2 stops	No info	13882	0	0
3	1 stop	No info	6218	0	0
4	1 stop	No info	13302	0	0

	Source_Chennai	Source_Delhi	Source_Kolkata	Source_Mumbai	\
0	0	0	0	0	
1	0	0	1	0	
2	0	1	0	0	
3	0	0	1	0	
4	0	0	0	0	

	Destination_Bangalore	Destination_Cochin	Destination_Delhi	\
0	0	0	0	
1	1	0	0	
2	0	1	0	
3	1	0	0	
4	0	0	0	

	Destination_Hyderabad	Destination_Kolkata	Destination_New Delhi
0	0	0	1
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	1

[5 rows x 28 columns]

Below, we will compare our original data frame with the transformed one.

```
[15]: data.shape
```

```
[15]: (10683, 11)
```

```
[16]: data1.shape
```

```
[16]: (10683, 28)
```

As we can see, we went from 11 original features in our dataset to 38. This is because *Pandas* `get_dummies()` approach when applied to a column with different categories (e.g. different airlines) will produce a new column (variable) for each unique categorical value (for each unique airline). It will place a one in the column corresponding to the categorical value present for that observation.

1.7 Exercise 2

In this exercise, use `value_counts()` to determine the values distribution of the 'Total_Stops' parameter.

```
[17]: # Enter your code and run the cell
data['Total_Stops'].value_counts()
```

```
[17]: 1 stop      5625
non-stop    3492
2 stops     1520
3 stops       45
4 stops        1
Name: Total_Stops, dtype: int64
```

Solution (Click Here)

    <code>

```
data["Total_Stops"].value_counts()
```

Label Encoding Since 'Total_Stops' is originally a categorical data type, we also need to convert it into numerical one. For this, we can perform a label encoding, where values are manually assigned to the corresponding keys, like "0" to a "non-stop", using the `replace()` function.

```
[18]: data1.replace({"non-stop":0,"1 stop":1,"2 stops":2,"3 stops":3,"4 stops":
↪4},inplace=True)
data1.head()
```

```
[18]:  Date_of_Journey      Route Dep_Time  Arrival_Time  Duration  \
0      24/03/2019      BLR → DEL    22:20    01:10 22 Mar    2h 50m
1      1/05/2019  CCU → IXR → BBI → BLR    05:50           13:15    7h 25m
2      9/06/2019  DEL → LKO → BOM → COK    09:25    04:25 10 Jun     19h
3     12/05/2019      CCU → NAG → BLR    18:05           23:30    5h 25m
4     01/03/2019      BLR → NAG → DEL    16:50           21:35    4h 45m

   Total_Stops  Additional_Info  Price  Airline_Air  Asia  Airline_Air  India  \
0            0         No info   3897           0      0           0
1            2         No info   7662           0           1
2            2         No info  13882           0           0
3            1         No info   6218           0           0
4            1         No info  13302           0           0

   ...  Source_Chennai  Source_Delhi  Source_Kolkata  Source_Mumbai  \
0  ...              0              0              0              0
1  ...              0              0              1              0
2  ...              0              1              0              0
3  ...              0              0              1              0
4  ...              0              0              0              0

   Destination_Banglore  Destination_Cochin  Destination_Delhi  \
0                    0                    0                    0
1                    1                    0                    0
```


2	0	1	0
3	1	0	0
4	0	0	0

	Destination_Hyderabad	Destination_Kolkata	Destination_New Delhi
0	0	0	1
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	1

[5 rows x 28 columns]

1.7.1 Date Time Transformations

Transforming the ‘Duration’ time column Here, we will take a closer look at the ‘Duration’ variable. Duration is the time taken by a plane to reach its destination. It is the difference between the ‘Dep_Time’ and ‘Arrival_Time’. In our dataset, the ‘Duration’ is expressed as a string, in hours and minutes. To be recognized by machine learning algorithms, we also need to transform it into numerical type.

The code below will iterate through each record in ‘Duration’ column and split it into hours and minutes, as two additional separate columns. Also, we want to add the ‘Duration_hours’ (in minutes) to the ‘Duration_minutes’ column to obtain a ‘Duration_Total_mins’ time, in minutes. The total duration time column will be useful feature for any regression type of analysis.

```
[19]: duration = list(data1['Duration'])
for i in range(len(duration)) :
    if len(duration[i].split()) != 2:
        if 'h' in duration[i] :
            duration[i] = duration[i].strip() + ' 0m'
        elif 'm' in duration[i] :
            duration[i] = '0h {}'.format(duration[i].strip())
dur_hours = []
dur_minutes = []

for i in range(len(duration)) :
    dur_hours.append(int(duration[i].split()[0][:-1]))
    dur_minutes.append(int(duration[i].split()[1][:-1]))

data1['Duration_hours'] = dur_hours
data1['Duration_minutes'] =dur_minutes
data1.loc[:, 'Duration_hours'] *= 60
data1['Duration_Total_mins']= data1['Duration_hours']+data1['Duration_minutes']
```

Print ‘data1’ data frame to see the newly created columns.

```
[20]: data1.head()
```

```
[20]:  Date_of_Journey      Route Dep_Time  Arrival_Time Duration \
0      24/03/2019      BLR → DEL    22:20    01:10 22 Mar    2h 50m
1      1/05/2019  CCU → IXR → BBI → BLR    05:50          13:15    7h 25m
2      9/06/2019  DEL → LKO → BOM → COK    09:25    04:25 10 Jun      19h
3      12/05/2019      CCU → NAG → BLR    18:05          23:30    5h 25m
4      01/03/2019      BLR → NAG → DEL    16:50          21:35    4h 45m

      Total_Stops Additional_Info  Price  Airline_Air Asia  Airline_Air India \
0              0          No info   3897              0              0
1              2          No info   7662              0              1
2              2          No info  13882              0              0
3              1          No info   6218              0              0
4              1          No info  13302              0              0

      ... Source_Mumbai Destination_Banglore Destination_Cochin \
0      ...              0              0              0
1      ...              0              1              0
2      ...              0              0              1
3      ...              0              1              0
4      ...              0              0              0

      Destination_Delhi Destination_Hyderabad Destination_Kolkata \
0              0              0              0
1              0              0              0
2              0              0              0
3              0              0              0
4              0              0              0

      Destination_New Delhi  Duration_hours  Duration_minutes \
0              1          120          50
1              0          420          25
2              0         1140           0
3              0          300          25
4              1          240          45

      Duration_Total_mins
0              170
1              445
2             1140
3              325
4              285
```

```
[5 rows x 31 columns]
```

As you have noticed, three new columns were created: 'Duration_hours', 'Duration_minutes', and

‘Duration_Total_mins’ - all numerical values.

Transforming the ‘Departure’ and ‘Arrival’ Time Columns Now, we will transform the ‘Dep_Time’ and ‘Arrival_Time’ columns to the appropriate date and time format. We will use *pandas* `to_datetime()` function for this.

We will split the ‘Dep_Time’ and ‘Arrival_Time’ columns into their corresponding hours and minutes columns.

```
[21]: data1["Dep_Hour"] = pd.to_datetime(data1['Dep_Time']).dt.hour
      data1["Dep_Min"] = pd.to_datetime(data1['Dep_Time']).dt.minute
```

1.8 Exercise 3

Now, let’s transform the ‘Arrival_Time’ column.

```
[24]: # Enter your code and run the cell
      data1['Arrival_Hour'] = pd.to_datetime(data1['Arrival_Time']).dt.hour
      data1['Arrival_Min'] = pd.to_datetime(data1['Arrival_Time']).dt.minute
```

Solution (Click Here)

    <code>

```
data1["Arrival_Hour"] = pd.to_datetime(data1['Arrival_Time']).dt.hour
data1["Arrival_Min"] = pd.to_datetime(data1['Arrival_Time']).dt.minute
```

Splitting ‘Departure/Arrival_Time’ into Time Zones To further transform our ‘Departure/Arrival_Time’ column, we can break down the 24 hours format for the departure and arrival time into 4 different time zones: night, morning, afternoon, and evening. This might be an interesting feature engineering technique to see what time of a day has the most arrivals/departures.

One way to do this transformation is by using *pandas* `cut()` function.

```
[23]: data1['dep_timezone'] = pd.cut(data1.Dep_Hour, [0,6,12,18,24],
      ↪ labels=['Night', 'Morning', 'Afternoon', 'Evening'])
      data1['dep_timezone']
```

```
[23]: 0      Evening
      1      Night
      2      Morning
      3      Afternoon
      4      Afternoon
      ...
10678    Evening
10679    Evening
10680    Morning
10681    Morning
10682    Morning
      Name: dep_timezone, Length: 10683, dtype: category
```

```
Categories (4, object): ['Night' < 'Morning' < 'Afternoon' < 'Evening']
```

1.9 Exercise 4

Now, let's transform the 'Arrival_Time' column into its corresponding time zones, as shown in the example above.

```
[ ]: # Enter your code and run the cell
data1['Arrival_timezones'] = pd.cut(data1['Arrival_Hour'], [0,6,12,18,24],
    ↪labels=['Night', 'Morning', 'Afternoon', 'Evening'])
```

[Solution \(Click Here\)](#)

    <code>

```
data1["Arrival_Hour"] = pd.to_datetime(data1['Arrival_Time']).dt.hour
data1['arr_timezone'] = pd.cut(data1.Arrival_Hour, [0,6,12,18,24], labels=['Night', 'Morning', 'Afternoon', 'Evening'])
```

Transforming the 'Date_of_Journey' Column Similar to the departure/arrival time, we will now extract some information from the 'date_of_journey' column, which is also an object type and can not be used for any machine learning algorithm yet.

So, we will extract the month information first and store it under the 'Month' column name.

```
[25]: data1['Month'] = pd.to_datetime(data1["Date_of_Journey"], format="%d/%m/%Y").dt.
    ↪month
```

1.10 Exercise 5

Now, let's create 'Day' and 'Year' columns in a similar way.

```
[26]: # Enter your code and run the cell
data1['Day'] = pd.to_datetime(data1["Date_of_Journey"], format="%d/%m/%Y").dt.day
data1['Year'] = pd.to_datetime(data1["Date_of_Journey"], format="%d/%m/%Y").dt.
    ↪year
```

[Solution \(Click Here\)](#)

    <code>

```
data1['Day'] = pd.to_datetime(data1["Date_of_Journey"], format="%d/%m/%Y").dt.day
data1['Year'] = pd.to_datetime(data1["Date_of_Journey"], format="%d/%m/%Y").dt.year
```

Additionally, we can extract the day of the week name by using `dt.day_name()` function.

```
[27]: data1['day_of_week'] = pd.to_datetime(data1['Date_of_Journey']).dt.day_name()
```

1.11 Feature Selection

Here, we will select only those attributes which best explain the relationship of the independent variables with respect to the target variable, 'price'. There are many methods for feature selection,

building the heatmap and calculating the correlation coefficients scores are the most commonly used ones.

First, we will select only the relevant and newly transformed variables (and exclude variables such as 'Route', 'Additional_Info', and all the original categorical variables), and place them into a 'new_data' data frame.

We will print all of our data1 columns.

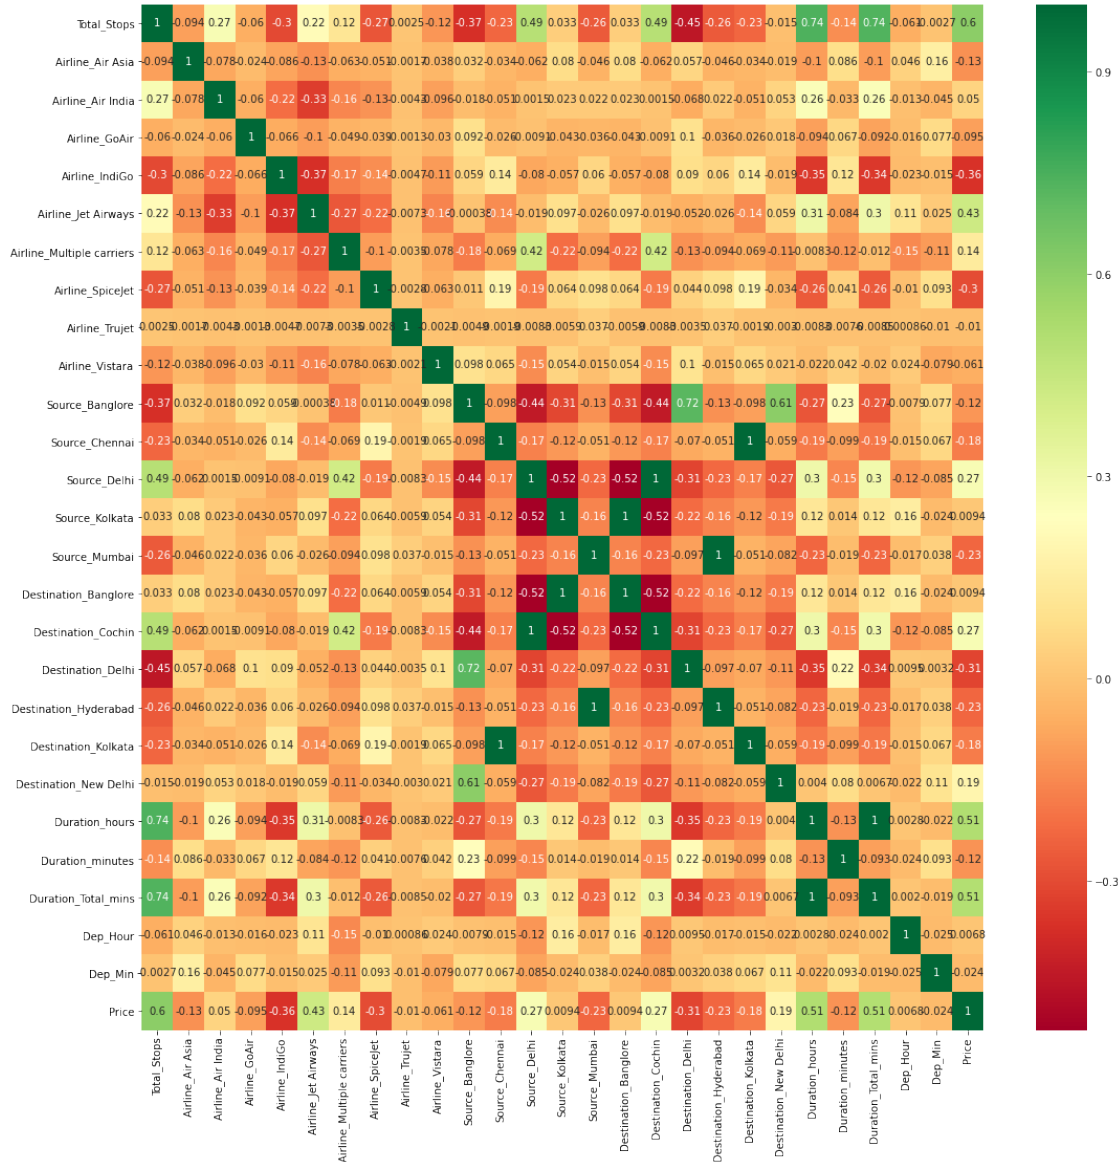
```
[28]: data1.columns
```

```
[28]: Index(['Date_of_Journey', 'Route', 'Dep_Time', 'Arrival_Time', 'Duration',  
        'Total_Stops', 'Additional_Info', 'Price', 'Airline_Air Asia',  
        'Airline_Air India', 'Airline_GoAir', 'Airline_IndiGo',  
        'Airline_Jet Airways', 'Airline_Multiple carriers', 'Airline_SpiceJet',  
        'Airline_Trujet', 'Airline_Vistara', 'Source_Banglore',  
        'Source_Chennai', 'Source_Delhi', 'Source_Kolkata', 'Source_Mumbai',  
        'Destination_Banglore', 'Destination_Cochin', 'Destination_Delhi',  
        'Destination_Hyderabad', 'Destination_Kolkata', 'Destination_New Delhi',  
        'Duration_hours', 'Duration_minutes', 'Duration_Total_mins', 'Dep_Hour',  
        'Dep_Min', 'dep_timezone', 'Arrival_Hour', 'Arrival_Min', 'Month',  
        'Day', 'Year', 'day_of_week'],  
        dtype='object')
```

```
[29]: new_data = data1.loc[:,['Total_Stops', 'Airline_Air Asia',  
        'Airline_Air India', 'Airline_GoAir', 'Airline_IndiGo',  
        'Airline_Jet Airways', 'Airline_Multiple carriers', 'Airline_SpiceJet',  
        'Airline_Trujet', 'Airline_Vistara', 'Source_Banglore',  
        'Source_Chennai', 'Source_Delhi', 'Source_Kolkata', 'Source_Mumbai',  
        'Destination_Banglore', 'Destination_Cochin', 'Destination_Delhi',  
        'Destination_Hyderabad', 'Destination_Kolkata', 'Destination_New Delhi',  
        'Duration_hours', 'Duration_minutes', 'Duration_Total_mins', 'Dep_Hour',  
        'Dep_Min', 'dep_timezone', 'Price']]
```

Now we will construct a `heatmap()`, using the *seaborn* library with a newly formed data frame, 'new_data'.

```
[30]: plt.figure(figsize=(18,18))  
sns.heatmap(new_data.corr(),annot=True,cmap='RdYlGn')  
  
plt.show()
```



From the heatmap above, extreme green means highly positively correlated features (relationship between two variables in which both variables move in the same direction), extreme red means negatively correlated features (relationship between two variables in which an increase in one variable is associated with a decrease in the other).

Now, we can use the `corr()` function to calculate and list the correlation between all independent variables and the 'price'.


```
[32]: features = new_data.corr()['Price'].sort_values(ascending= False)
      features
```

```
[32]: Price                1.000000
      Total_Stops          0.603891
      Duration_hours       0.508672
      Duration_Total_mins  0.506371
      Airline_Jet Airways  0.428490
      Source_Delhi         0.270619
      Destination_Cochin   0.270619
      Destination_New Delhi 0.189785
      Airline_Multiple carriers 0.141087
      Airline_Air India    0.050346
      Destination_Bangalore 0.009377
      Source_Kolkata       0.009377
      Dep_Hour             0.006819
      Airline_Trujet      -0.010380
      Dep_Min              -0.024492
      Airline_Vistara      -0.060503
      Airline_GoAir        -0.095146
      Source_Bangalore     -0.118026
      Duration_minutes     -0.124874
      Airline_Air Asia     -0.133044
      Source_Chennai       -0.179216
      Destination_Kolkata  -0.179216
      Destination_Hyderabad -0.230745
      Source_Mumbai        -0.230745
      Airline_SpiceJet      -0.296552
      Destination_Delhi    -0.313401
      Airline_IndiGo       -0.361048
      Name: Price, dtype: float64
```

We can also plot these correlation coefficients for easier visualization.

```
[33]: features.plot(kind='bar',figsize=(10,8))
```

```
[33]: <AxesSubplot:>
```



From the graph above, we can deduct some of the highly correlated features and select only those ones for any future analysis.

1.12 Feature Extraction using Principal Component Analysis (Optional)

1.12.1 PCA with Scikit-Learn

Dimentionality reduction is part of the feature extraction process that combines the existing features to produce more useful ones. The goal of dimensionality reduction is to simplify the data without losing too much information. Principal Component Analysis (PCA) is one of the most popular dimensionality reduction algorithms. First, it identifies the hyperplane that lies closest to the data, and then it projects the data onto it. In this way, a few multidimensional features are merged into one.

In the following portion of the lab, we will use `scikit-learn` library to perform some PCA on our data. To learn more about `scikit-learn` PCA, please visit this [documentation](#).

First, we must scale our data using the `StandardScaler()` function. We will assign all the independent variables to `x`, and the dependent variable, 'price', to `y`.

```
[34]: x = data1.loc[:,['Total_Stops', 'Airline_Air Asia',  
    'Airline_Air India', 'Airline_GoAir', 'Airline_IndiGo',  
    'Airline_Jet Airways', 'Airline_Multiple carriers', 'Airline_SpiceJet',  
    'Airline_Trujet', 'Airline_Vistara', 'Source_Bangalore',  
    'Source_Chennai', 'Source_Delhi', 'Source_Kolkata', 'Source_Mumbai',  
    'Destination_Bangalore', 'Destination_Cochin', 'Destination_Delhi',  
    'Destination_Hyderabad', 'Destination_Kolkata', 'Destination_New Delhi',  
    'Duration_hours', 'Duration_minutes', 'Duration_Total_mins', 'Dep_Hour',  
    'Dep_Min']]
```

```
[35]: y= data1.Price
```

```
[36]: scaler = StandardScaler()  
x=scaler.fit_transform(x.astype(np.float64))  
x
```

```
[36]: array([[ -1.22052384, -0.17544122, -0.44291155, ..., -0.93158255,  
    1.65425948, -0.23505036],  
    [ 1.74150619, -0.17544122,  2.25778713, ..., -0.39007152,  
   -1.30309491,  1.36349161],  
    [ 1.74150619, -0.17544122, -0.44291155, ...,  0.97847452,  
   -0.60724682,  0.0313733 ],  
    ...,  
    [ -1.22052384, -0.17544122, -0.44291155, ..., -0.91189124,  
   -0.78120884, -0.23505036],  
    [ -1.22052384, -0.17544122, -0.44291155, ..., -0.95127386,  
   -0.25932278,  0.29779696],  
    [ 1.74150619, -0.17544122,  2.25778713, ..., -0.28176932,  
   -0.4332848 ,  1.62991527]])
```

Once the data is scaled, we can apply the `fit_transform()` function to reduce the dimensionality of the dataset down to two dimensions.

```
[37]: pca = PCA(n_components = 2)  
pca.fit_transform(x)
```

```
[37]: array([[ -2.87557371, -0.55522498],  
    [ 0.31881914,  2.39238454],  
    [ 3.05930533, -0.5268359 ],  
    ...,  
    [ -2.2475578 , -0.58846432],  
    [ -2.69896781, -0.28882268],  
    [ 1.92548332, -1.10414924]])
```

1.12.2 Explained Variance Ratio

Another useful piece of information in PCA is the explained variance ratio of each principal component, available via the `explained_variance_ratio_` function. The ratio indicates the proportion of the dataset's variance that lies along each principal component. Let's look at the explained variance ratio of each of our two components.

```
[38]: explained_variance=pca.explained_variance_ratio_  
      explained_variance
```

```
[38]: array([0.17545521, 0.12110719])
```

The first component constitutes 17.54% of the variance and second component constitutes 12.11% of the variance between the features.

1.13 Exercise 6 (Optional)

In this exercise, experiment with the number of components to see how many dimensions our dataset could be reduced to in order to explain most of the variability between the features. Additionally, you can plot the components using bar plot to see how much variability each component represents.

```
[43]: # Enter your code and run the cell  
dim = list(range(len(x)))  
for d in dim:  
    ex_6 = PCA(n_components = d)  
    ex_6.fit_transform(x)  
  
    explained_variance=ex_6.explained_variance_ratio_  
    print(d, explained_variance)
```

```
0 []  
1 [0.17545521]  
2 [0.17545521 0.12110718]  
3 [0.17545514 0.12110711 0.09264443]  
4 [0.1754552 0.12110701 0.09264705 0.08279936]  
5 [0.17545521 0.12110699 0.09264918 0.08279969 0.06739467]  
6 [0.17545521 0.12110713 0.0926492 0.0828011 0.06739565 0.05274453]  
7 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819541]  
8 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819544 0.04491853]  
9 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819544 0.04491853 0.04123197]  
10 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819544 0.04491853 0.04123197 0.03950738]  
11 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819544 0.04491853 0.04123197 0.03950738 0.03861142]  
12 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645  
 0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285]
```

13 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233]

14 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233 0.03376868]

15 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233 0.03376868 0.02884552]

16 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233 0.03376868 0.02884552 0.02685501]

17 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233 0.03376868 0.02884552 0.02685501 0.02086503]

18 [0.17545521 0.12110719 0.0926492 0.08280111 0.06739565 0.05275645
0.04819544 0.04491853 0.04123197 0.03950738 0.03861142 0.03842285
0.03692233 0.03376868 0.02884552 0.02685501 0.02086503 0.00969103]

19 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 9.62517115e-34]

20 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 9.62517115e-34 9.62517115e-34]

21 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32]

22 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32 1.26344797e-32]

23 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32 1.26344797e-32 4.82937991e-33]

24 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02

```

4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32 1.26344797e-32 4.82937991e-33 2.51126109e-33]
25 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32 1.26344797e-32 4.82937991e-33 2.51126109e-33
1.19656174e-33]
26 [1.75455212e-01 1.21107185e-01 9.26492032e-02 8.28011061e-02
6.73956512e-02 5.27564550e-02 4.81954443e-02 4.49185343e-02
4.12319717e-02 3.95073767e-02 3.86114153e-02 3.84228457e-02
3.69223293e-02 3.37686807e-02 2.88455153e-02 2.68550107e-02
2.08650342e-02 9.69102925e-03 6.32050506e-32 5.33125976e-32
2.28515894e-32 1.26344797e-32 4.82937991e-33 2.51126109e-33
1.19656174e-33 2.04050468e-34]

```

```

-----
ValueError                                Traceback (most recent call last)
/tmp/ipykernel_721/2423832777.py in <module>
      3 for d in dim:
      4     ex_6 = PCA(n_components = d)
----> 5     ex_6.fit_transform(x)
      6
      7     explained_variance=ex_6.explained_variance_ratio_

~/conda/envs/python/lib/python3.7/site-packages/sklearn/decomposition/pca.py in
fit_transform(self, X, y)
    357
    358     """
--> 359     U, S, V = self._fit(X)
    360     U = U[:, :self.n_components_]
    361

~/conda/envs/python/lib/python3.7/site-packages/sklearn/decomposition/pca.py in
_fit(self, X)
    404     # Call different fits for either full or truncated SVD
    405     if self._fit_svd_solver == 'full':
--> 406         return self._fit_full(X, n_components)
    407     elif self._fit_svd_solver in ['arpack', 'randomized']:
    408         return self._fit_truncated(X, n_components, self.
_fit_svd_solver)

~/conda/envs/python/lib/python3.7/site-packages/sklearn/decomposition/pca.py in
_fit_full(self, X, n_components)

```

```

423                                     "min(n_samples, n_features)=%r with "
424                                     "svd_solver='full'"
--> 425                                     % (n_components, min(n_samples,
↳n_features)))
426         elif n_components >= 1:
427             if not isinstance(n_components, (numbers.Integral, np.
↳integer)):
ValueError: n_components=27 must be between 0 and min(n_samples, n_features)=26,
↳with svd_solver='full'

```

[]: *# Enter your code and run the cell*

Solution_part1 (Click Here)

    <code>

```
pca = PCA(n_components=7)
pca.fit_transform(x)
explained_variance=pca.explained_variance_ratio_
explained_variance
```

Solution_part2 (Click Here)

    <code>

```

with plt.style.context('dark_background'): plt.figure(figsize=(6, 4))
plt.bar(range(7), explained_variance, alpha=0.5, align='center',
        label='individual explained variance')
plt.ylabel('Explained variance ratio')
plt.xlabel('Principal components')
plt.legend(loc='best')
plt.tight_layout()

```

1.13.1 Choosing the Right Number of Dimensions

Instead of arbitrary choosing the number of dimensions to reduce down to, it is simpler to choose the number of dimensions that add up to a sufficiently large proportion of the variance, let's say 95%.

The following code performs PCA without reducing dimensionality, then computes the minimum number of dimensions required to preserve 95% of the variance.

```

[44]: pca = PCA()
      pca.fit(x)
      cumsum = np.cumsum(pca.explained_variance_ratio_)
      d = np.argmax(cumsum >=0.95) + 1

```

```
[45]: d
```


```
[45]: 16
```

There are 16 components required to meet 95% variance. Therefore, we could set `n_components = 16` and run PCA again. However, there is better way, instead of specifying the number of principal components you want to preserve, you can set `n_components` to be a float between 0.0 and 1.0, indicating the ratio of variance you wish to preserve.

```
[48]: pca = PCA(n_components=0.95)
      x_reduced = pca.fit_transform(x)
```

There is also a graphical way to determine the number of principal components in your analysis. It is to plot the explained variance as a function of the number of dimensions. There will usually be an elbow in the curve, where the explained variance stops growing fast. That point is usually the optimal point for the number of principal components.

```
[47]: px.area(
      x=range(1, cumsum.shape[0] + 1),
      y=cumsum,
      labels={"x": "# Components", "y": "Explained Variance"}
    )
```



2 Congratulations! - You have completed the lab

2.1 Author

[Svitlana Kramar](#)

2.2 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2022-01-17	0.1	Svitlana	Modified multiple areas

Copyright © 2020 IBM Corporation. All rights reserved.