

Stacking For Classification with Python

May 18, 2022

1 Stacking For Classification with Python

Estimated time needed: **45** minutes

1.1 Objectives

After completing this lab you will be able to:

- **Understand** what Stacking is and how it works
- **Understand** that Random Forests have less Correlation between predictors in their ensemble, improving accuracy
- **Apply** Stacking
- **Understand** Hyperparameters selection in Stacking

Stacking takes several classification models called base learners and uses their output as the input for the meta-classifier. Consider the figure below the base learners $h_1(x)$, $h_2(x)$, $h_3(x)$, and $h_4(x)$ has the output $\hat{y} * 1$, $\hat{y} * 2$, $\hat{y} * 3$, $\hat{y} * 4$. These are used as an input to the meta classifier $H(\hat{y} * 1, \hat{y} * 2, \hat{y} * 3, \hat{y} * 4)$, makes the final prediction $\hat{y} = H(\hat{y} * 1, \hat{y} * 2, \hat{y} * 3, \hat{y} * 4)$.

Fig. 1 Stacking takes several classification models called base learners and uses their output as the input for the meta-classifier.

We can train all the models using all the data but this causes over-fitting. To get a better idea of how the algorithm works we use K-fold Cross-validation. We use K-1 folds to train the base classifiers and the last fold to train the meta classifier. We repeat the process using different combinations of each fold. This is shown in Fig 2 where the color-coded square represents the different runs and folds. Each row represents a different run of K fold cross-validation, each column is one of K folds where K=3. For each column, we use the same color columns to train the classifiers and the different color is used to train the meta classifier.

Fig. 2 K-fold Cross-validation to train Stacking classifier.

Table of contents

<a>Apply Staking Using Wine Data

<a>Cancer Data Example

First, let's install and import the required libraries:

```
[1]: # All Libraries required for this lab are listed below. The libraries,
      ↪pre-installed on Skills Network Labs are commented.
      # !mamba install -qy pandas==1.3.3 numpy==1.21.2 ipywidgets==7.4.2 scipy==7.4.2
      ↪tqdm==4.62.3 matplotlib==3.5.0 seaborn==0.9.0

      # You will need scikit-learn>=0.22.0 as StackingClassifier does not exist in
      ↪version <0.22.0
      !mamba install -c conda-forge -qy scikit-learn=0.22.0

      # Note: If your environment doesn't support "!mamba install", use "!pip install"
```

Package Size	Version	Build	Channel
-----------------	---------	-------	---------

Install:

+ atk-1.0	2.36.0	h3371d22_4	
conda-forge/linux-64 574kB			
+ cached-property	1.5.2	hd8ed1ab_1	
conda-forge/noarch 4kB			
+ cached_property	1.5.2	pyha770c72_1	
conda-forge/noarch 11kB			
+ font-ttf-dejavu-sans-mono	2.37	hab24e00_0	
conda-forge/noarch 397kB			
+ font-ttf-inconsolata	3.000	h77eed37_0	
conda-forge/noarch 97kB			
+ font-ttf-source-code-pro	2.038	h77eed37_0	
conda-forge/noarch 701kB			
+ font-ttf-ubuntu	0.83	hab24e00_0	
conda-forge/noarch 2MB			
+ fonts-conda-ecosystem	1	0	
conda-forge/noarch 4kB			
+ fonts-conda-forge	1	0	
conda-forge/noarch 4kB			
+ gdk-pixbuf	2.42.8	hff1cb4f_0	
conda-forge/linux-64 609kB			
+ gtk2	2.24.33	h90689f9_2	
conda-forge/linux-64 8MB			
+ gts	0.7.6	h64030ff_2	
conda-forge/linux-64 421kB			
+ jbig	2.1	h7f98852_2003	
conda-forge/linux-64 44kB			
+ joblib	1.1.0	pyhd8ed1ab_0	
conda-forge/noarch 215kB			

+ lerc		3.0	h9c3ff4c_0
conda-forge/linux-64	222kB		
+ libdeflate		1.10	h7f98852_0
conda-forge/linux-64	79kB		
+ libgd		2.3.3	h283352f_2
conda-forge/linux-64	278kB		
+ libllvm13		13.0.1	hf817b99_2
conda-forge/linux-64	35MB		
+ librsvg		2.52.5	h0a9e6e8_2
conda-forge/linux-64	6MB		
+ libwebp-base		1.2.2	h7f98852_1
conda-forge/linux-64	844kB		
+ scikit-learn		0.22	py37hcdab131_1
conda-forge/linux-64	7MB		

Change:

- cairo		1.16.0	h6cf1ce9_1008
installed			
+ cairo		1.16.0	ha12eb4b_1010
conda-forge/linux-64	2MB		
- krb5		1.19.3	h3790be6_0
installed			
+ krb5		1.19.3	h08a2579_0
conda-forge/linux-64	2MB		
- leptonica		1.78.0	h42ed529_2
installed			
+ leptonica		1.78.0	ha0c4403_3
conda-forge/linux-64	3MB		
- libevent		2.1.10	h9b69904_4
installed			
+ libevent		2.1.10	h28343ad_4
conda-forge/linux-64	1MB		
- libnghttp2		1.47.0	h727a467_0
installed			
+ libnghttp2		1.47.0	he49606f_0
conda-forge/linux-64	844kB		
- libssh2		1.10.0	ha56f1ee_2
installed			
+ libssh2		1.10.0	ha35d2d1_2
conda-forge/linux-64	238kB		
- libxml2		2.9.12	h72842e0_0
installed			
+ libxml2		2.9.12	h885dcf4_1
conda-forge/linux-64	778kB		
- python		3.7.12	hb7a2778_100_cpython

```

installed
+ python
conda-forge/linux-64      60MB      3.7.12  hf930737_100_cpython
- qt
installed
+ qt
conda-forge/linux-64      103MB     5.12.9  hda022c4_4
- tesseract
installed
+ tesseract
conda-forge/linux-64      325MB     4.1.1   h9862bf9_2
4.1.1   h58164bb_3

Upgrade:

- cryptography
installed
+ cryptography
pkgs/main/linux-64        1MB      36.0.2  py37h38fbfac_1
37.0.1  py37h9ce1e76_0
- ffmpeg
installed
+ ffmpeg
conda-forge/linux-64      10MB     4.1.3   h167e202_0
4.3.2   h37c90e5_3
- graphviz
installed
+ graphviz
conda-forge/linux-64      3MB      2.40.1  h0511662_2
3.0.0   h5abf519_1
- grpcio
installed
+ grpcio
conda-forge/linux-64      3MB      1.45.0  py37he500948_0
1.46.1  py37h0327239_0
- gst-plugins-base
installed
+ gst-plugins-base
conda-forge/linux-64      3MB      1.18.5  hf529b03_3
1.20.1  hcf0ee16_1
- gstreamer
installed
+ gstreamer
conda-forge/linux-64      2MB      1.18.5  h9f60fe5_3
1.20.2  hd4edc92_1
- h5py
installed
+ h5py
conda-forge/linux-64      1MB      2.10.0  nompi_py37h513d04c_102
3.6.0   nompi_py37hd308b1e_100
- harfbuzz
installed
+ harfbuzz
conda-forge/linux-64      2MB      2.9.1   h83ec7ef_1
3.4.0   hb4a5f5f_0

```

- hdf5		1.10.5	nompi_h5b725eb_1114
installed			
+ hdf5		1.12.1	nompi_h4df4325_104
conda-forge/linux-64	4MB		
- icu		68.2	h9c3ff4c_0
installed			
+ icu		69.1	h9c3ff4c_0
conda-forge/linux-64	14MB		
- libarchive		3.5.1	h3f442fb_1
installed			
+ libarchive		3.5.2	hed592e5_1
conda-forge/linux-64	2MB		
- libclang		11.1.0	default_ha53f305_1
installed			
+ libclang		13.0.1	default_hc23dcda_0
conda-forge/linux-64	12MB		
- libcurl		7.82.0	h7bff187_0
installed			
+ libcurl		7.83.1	h2283fc2_0
conda-forge/linux-64	352kB		
- libopencv		3.4.9	py37_2
installed			
+ libopencv		4.5.3	py37hb6cea29_4
conda-forge/linux-64	36MB		
- libpq		13.5	hd57d9b9_1
installed			
+ libpq		14.3	he2d8382_0
conda-forge/linux-64	3MB		
- libtiff		4.1.0	hc3755c2_3
installed			
+ libtiff		4.3.0	h542a066_3
conda-forge/linux-64	653kB		
- libwebp		1.0.2	h56121f0_5
installed			
+ libwebp		1.2.2	h3452ae3_0
conda-forge/linux-64	87kB		
- mysql-common		8.0.25	ha770c72_0
installed			
+ mysql-common		8.0.29	h26416b9_1
conda-forge/linux-64	2MB		
- mysql-libs		8.0.25	h935591d_0
installed			
+ mysql-libs		8.0.29	hbc51c84_1
conda-forge/linux-64	2MB		
- opencv		3.4.9	py37_2
installed			
+ opencv		4.5.3	py37h89c1867_4
conda-forge/linux-64	22kB		

```

-openh264                                1.8.0  hdbcaa40_1000
installed
+openh264                                2.1.1  h780b84a_0
conda-forge/linux-64                    2MB
-openssl                                1.1.1n  h166bdaf_0
installed
+openssl                                3.0.3  h166bdaf_0
conda-forge/linux-64                    3MB
-pango                                  1.42.4  h69149e4_5
installed
+pango                                  1.50.5  h4dcc4a0_1
conda-forge/linux-64                    466kB
-py-opencv                              3.4.9  py37h5ca1d4c_2
installed
+py-opencv                              4.5.3  py37h6531663_4
conda-forge/linux-64                    1MB
-x264                                  1!152.20180806  h14c3975_0
installed
+x264                                  1!161.3030  h7f98852_1
conda-forge/linux-64                    3MB
-zstd                                  1.4.9  ha95c52a_0
installed
+zstd                                  1.5.2  h8a70e8d_1
conda-forge/linux-64                    462kB

```

Downgrade:

```

-libprotobuf                            3.20.0  h6239696_0
installed
+libprotobuf                            3.18.1  h780b84a_0
conda-forge/linux-64                    3MB
-protobuf                              3.20.0  py37hd23a5d3_4
installed
+protobuf                              3.18.1  py37hcd2ae1e_0
conda-forge/linux-64                    353kB

```

Summary:

```

Install: 21 packages
Change: 10 packages
Upgrade: 26 packages
Downgrade: 2 packages

```

Total download: 673MB

Preparing transaction: ...working... done
Verifying transaction: ...working... done
Executing transaction: ...working...

done

```
[19]: import pandas as pd
      # import pylab as plt
      import numpy as np
      import scipy.optimize as opt
      from sklearn import preprocessing
      import matplotlib.pyplot as plt
      %matplotlib inline
      from sklearn import metrics
      from sklearn.model_selection import GridSearchCV
      import seaborn as sns
      from sklearn import preprocessing
      from sklearn.ensemble import StackingClassifier
      from sklearn.svm import SVC
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.linear_model import LogisticRegression
```

```
-----
AttributeError                                Traceback (most recent call last)
/tmp/ipykernel_1275/82920220.py in <module>
     10 import seaborn as sns
     11 from sklearn import preprocessing
----> 12 from sklearn.ensemble import StackingClassifier
     13 from sklearn.svm import SVC
     14 from sklearn.neighbors import KNeighborsClassifier

~/conda/envs/python/lib/python3.7/site-packages/sklearn/ensemble/__init__.py in
<=> <module>
      5
      6 from ._base import BaseEnsemble
----> 7 from ._forest import RandomForestClassifier
      8 from ._forest import RandomForestRegressor
      9 from ._forest import RandomTreesEmbedding

~/conda/envs/python/lib/python3.7/site-packages/sklearn/ensemble/_forest.py in
<=> <module>
     54 from ..metrics import r2_score
     55 from ..preprocessing import OneHotEncoder
```

```

---> 56 from ..tree import (DecisionTreeClassifier, DecisionTreeRegressor,
57                          ExtraTreeClassifier, ExtraTreeRegressor)
58 from ..tree._tree import DTYPE, DOUBLE

~/conda/envs/python/lib/python3.7/site-packages/sklearn/tree/__init__.py in
↳<module>
    4 """
    5
----> 6 from ._classes import BaseDecisionTree
    7 from ._classes import DecisionTreeClassifier
    8 from ._classes import DecisionTreeRegressor

~/conda/envs/python/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳<module>
    38 from ..utils.validation import check_is_fitted
    39
---> 40 from ._criterion import Criterion
    41 from ._splitter import Splitter
    42 from ._tree import DepthFirstTreeBuilder

~/conda/envs/python/lib/python3.7/site-packages/sklearn/tree/_criterion.
↳cpython-37m-x86_64-linux-gnu.so in init sklearn.tree._criterion()

AttributeError: type object 'sklearn.tree._criterion.array' has no attribute
↳'__reduce_cython__'

```

Ignore error warnings

```

[3]: import warnings
      warnings.filterwarnings('ignore')

```

This function will calculate the accuracy of the training and testing data given a model.

```

[4]: def get_accuracy(X_train, X_test, y_train, y_test, model):
      return {"test Accuracy": metrics.accuracy_score(y_test, model.
↳predict(X_test)), "trian Accuracy": metrics.accuracy_score(y_train, model.
↳predict(X_train))}

```

Apply Staking Using Wine Data

The class is an essential factor in determining the quality of the wine; this dataset uses chemical analysis of wines grown in the same region in Italy from three different cultivars. Your task is to determine the class of the wine using the features from the chemical analysis. The features or attributes include

For more info here ,let's load the dataset:


```
[5]: df = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-ML241EN-SkillsNetwork/labs/datasets/wine.data", names= ['Class', 'Alcohol', 'Malic acid', 'Ash', 'Alcalinity of ash', 'Magnesium', 'Total phenols', 'Flavanoids', 'Nonflavanoid phenols', 'Proanthocyanins', 'Color intensity', 'Hue', 'OD280/OD315 of diluted wines', 'Proline'])

df.head()
```

```
[5]:
```

	Class	Alcohol	Malic acid	Ash	Alcalinity of ash	Magnesium \
0	1	14.23	1.71	2.43	15.6	127
1	1	13.20	1.78	2.14	11.2	100
2	1	13.16	2.36	2.67	18.6	101
3	1	14.37	1.95	2.50	16.8	113
4	1	13.24	2.59	2.87	21.0	118

	Total phenols	Flavanoids	Nonflavanoid phenols	Proanthocyanins \
0	2.80	3.06	0.28	2.29
1	2.65	2.76	0.26	1.28
2	2.80	3.24	0.30	2.81
3	3.85	3.49	0.24	2.18
4	2.80	2.69	0.39	1.82

	Color intensity	Hue	OD280/OD315 of diluted wines	Proline
0	5.64	1.04	3.92	1065
1	4.38	1.05	3.40	1050
2	5.68	1.03	3.17	1185
3	7.80	0.86	3.45	1480
4	4.32	1.04	2.93	735

We see all the dataset is comprised of numerical values using the method dtypes

```
[6]: df.dtypes
```

```
[6]: Class          int64
Alcohol         float64
Malic acid      float64
Ash            float64
Alcalinity of ash float64
Magnesium       int64
Total phenols   float64
Flavanoids      float64
Nonflavanoid phenols float64
Proanthocyanins float64
Color intensity float64
Hue            float64
```

```
OD280/OD315 of diluted wines    float64
Proline                          int64
dtype: object
```


the column class has the class of the wine, we can use the method `unique()` to obtain the classes:

```
[7]: df['Class'].unique()
```

```
[7]: array([1, 2, 3])
```

We can examine the correlation between each feature and the class variable. By examining the first row or column we see the features are correlated with the class variable.

```
[8]: colormap = plt.cm.RdBu
plt.figure(figsize=(14,12))
plt.title('Pearson Correlation of Features', y=1.05, size=15)
sns.heatmap(df.astype(float).corr(),linewidths=0.1,vmax=1.0,
            square=True, cmap=colormap, linecolor='white', annot=True)
plt.show()
```



We can also examine the Pairplot between pairs of features and the histogram; color-coded to each class. We see the separation between different classes:

```
[9]: # May need to specify bandwidth (bw) in order to plot, else can ignore.
sns.pairplot(df, hue="Class") #, diag_kws={'bw': 0.2})
```

```
-----
RuntimeError                                Traceback (most recent call last)
/tmp/ipykernel_1275/2512947022.py in <module>
      1 # May need to specify bandwidth (bw) in order to plot, else can ignore.
----> 2 sns.pairplot(df, hue="Class") #, diag_kws={'bw': 0.2})

~/conda/envs/python/lib/python3.7/site-packages/seaborn/axisgrid.py in
pairplot(data, hue, hue_order, palette, vars, x_vars, y_vars, kind, diag_kind,
markers, height, aspect, dropna, plot_kws, diag_kws, grid_kws, size)
```

```

2109         diag_kws.setdefault("shade", True)
2110         diag_kws["legend"] = False
-> 2111         grid.map_diag(kdeplot, **diag_kws)
2112
2113         # Maybe plot on the off-diagonals

~/conda/envs/python/lib/python3.7/site-packages/seaborn/axisgrid.py in
↳ map_diag(self, func, **kwargs)
1397         color = fixed_color
1398
-> 1399         func(data_k, label=label_k, color=color, **kwargs)
1400
1401         self._clean_axis(ax)

~/conda/envs/python/lib/python3.7/site-packages/seaborn/distributions.py in
↳ kdeplot(data, data2, shade, vertical, kernel, bw, gridsize, cut, clip, legend,
↳ cumulative, shade_lowest, cbar, cbar_ax, cbar_kws, ax, **kwargs)
689         ax = _univariate_kdeplot(data, shade, vertical, kernel, bw,
690                                   gridsize, cut, clip, legend, ax,
--> 691                                   cumulative=cumulative, **kwargs)
692
693         return ax

~/conda/envs/python/lib/python3.7/site-packages/seaborn/distributions.py in
↳ _univariate_kdeplot(data, shade, vertical, kernel, bw, gridsize, cut, clip,
↳ legend, ax, cumulative, **kwargs)
281         x, y = _statsmodels_univariate_kde(data, kernel, bw,
282                                              gridsize, cut, clip,
--> 283                                              cumulative=cumulative)
284
285         else:
286             # Fall back to scipy if missing statsmodels

~/conda/envs/python/lib/python3.7/site-packages/seaborn/distributions.py in
↳ _statsmodels_univariate_kde(data, kernel, bw, gridsize, cut, clip, cumulative)
353         fft = kernel == "gau"
354         kde = smnp.KDEUnivariate(data)
--> 355         kde.fit(kernel, bw, fft, gridsize=gridsize, cut=cut, clip=clip)
356         if cumulative:
357             grid, y = kde.support, kde.cdf


~/conda/envs/python/lib/python3.7/site-packages/statsmodels/nonparametric/kde.p
↳ in fit(self, kernel, bw, fft, weights, gridsize, adjust, cut, clip)
173         gridsize=gridsize,
174         clip=clip,
--> 175         cut=cut,
176     )
177     else:

```

```
~/conda/envs/python/lib/python3.7/site-packages/statsmodels/nonparametric/kde.p
↳ in kdensityfft(x, kernel, bw, weights, gridsize, adjust, clip, cut, retgrid)
    555     elif isinstance(bw, str):
    556         # if bw is None, select optimal bandwidth for kernel
--> 557         bw = bandwidths.select_bandwidth(x, bw, kern)
    558         # will cross-val fit this pattern?
    559     else:

~/conda/envs/python/lib/python3.7/site-packages/statsmodels/nonparametric/
↳ bandwidths.py in select_bandwidth(x, bw, kernel)
    180         "Either provide the bandwidth during initialization or use
↳ " \
    181         "an alternative method."
--> 182         raise RuntimeError(err)
    183     else:
    184         return bandwidth
```

RuntimeError: Selected KDE bandwidth is 0. Cannot estimate density. Either
↳ provide the bandwidth during initialization or use an alternative method.



1.1.1 Data Pre-Processing and Selection

Let's examine the feature list

```
[11]: features=list(df)
      features[1:]
```

```
[11]: ['Alcohol',
      'Malic acid',
      'Ash',
      'Alcalinity of ash',
      'Magnesium',
      'Total phenols',
```

```
'Flavanoids',
'Nonflavanoid phenols',
'Proanthocyanins',
'Color intensity',
'Hue',
'OD280/OD315 of diluted wines',
'Proline']
```

We assign the class variables to y and feature variables to X

```
[12]: y,X=df[features[0]] ,df[features[1:]]
```

We can standardize the data

```
[13]: scaler = preprocessing.StandardScaler().fit(X)
X= scaler.transform(X)
```

We can check if the data is standardized by checking the mean and standard deviation, which are approximately zero:

```
[14]: X.mean(axis=0)
```

```
[14]: array([-8.38280756e-16, -1.19754394e-16, -8.37033314e-16, -3.99181312e-17,
          -3.99181312e-17,  0.00000000e+00, -3.99181312e-16,  3.59263181e-16,
          -1.19754394e-16,  2.49488320e-17,  1.99590656e-16,  3.19345050e-16,
          -1.59672525e-16])
```

```
[15]: X.std(axis=0)
```

```
[15]: array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.])
```

In Random Forest, we would use these data subsets to train each node of a tree.

1.1.2 Train/Test dataset

We split our dataset into train and test set:

```
[16]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.3,
    random_state=1)
print ('Train set', X_train.shape, y_train.shape)
print ('Test set', X_test.shape, y_test.shape)
```

Train set (124, 13) (124,)

Test set (54, 13) (54,)

Stacking consists of creating a Stacking Classifier object, but first, you require a dictionary of estimators (individual model objects or base learners). The key of the dictionary is a name that is up to you, we use the usual acronym for the model. The value is the model object in this case SVC for Support Vector Classifier, dt for Decision Tree Classifier and KNN for K Neighbors Classifier.

```
[17]: estimators =_
      ↪[('SVM',SVC(random_state=42)),('KNN',KNeighborsClassifier()),('dt',DecisionTreeClassifier())
      estimators
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_1275/196601360.py in <module>
----> 1 estimators =_
      ↪[('SVM',SVC(random_state=42)),('KNN',KNeighborsClassifier()),('dt',DecisionTreeClassifier())
      2 estimators

NameError: name 'SVC' is not defined
```

To train the final model we create a Stacking Classifier, this combines the base estimators using the meta estimator. The meta-classifier is determined by the parameter `final_estimator` in this case we use Logistic Regression, we also input the base classifiers using the `estimators` parameter and fit the model.

```
[20]: clf = StackingClassifier( estimators=estimators, final_estimator=_
      ↪LogisticRegression())
      clf.fit(X_train, y_train)
      clf
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_1275/2258820003.py in <module>
----> 1 clf = StackingClassifier( estimators=estimators, final_estimator=_
      ↪LogisticRegression())
      2 clf.fit(X_train, y_train)
      3 clf

NameError: name 'StackingClassifier' is not defined
```

We can make a prediction

```
[ ]: yhat=clf.predict(X_test)
      yhat
```

We can obtain the training and testing accuracy, we see the model performs well.

```
[ ]: get_accuracy(X_train, X_test, y_train, y_test, clf)
```

Note: Like most complex models Stacking is prone to overfitting

Practice

Create a Stacking Classifier object as before but exchange the Decision Tree Classifier with the SVM classifier. Calculate the accuracy on the training and testing data.

[]:

[Click here for the solution](#)

```
estimators = [('SVM', SVC(random_state=42)), ('KNN', KNeighborsClassifier()), ('lr', LogisticRegression())]
clf = StackingClassifier(estimators=estimators, final_estimator= DecisionTreeClassifier())
clf.fit(X_train, y_train)
```

```
get_accuracy(X_train, X_test, y_train, y_test, clf)
```

GridSearchCV and Stacking Classifiers

Imagine that you are a medical researcher compiling data for a study. You have collected data about a set of patients, all of whom suffered from the same illness. During their course of treatment, each patient responded to one of 5 medications, Drug A, Drug B, Drug c, Drug x and y.

Part of your job is to build a model to find out which drug might be appropriate for a future patient with the same illness. The features of this dataset are Age, Sex, Blood Pressure, and the Cholesterol of the patients, and the target is the drug that each patient responded to.

It is a sample of multiclass classifier, and you can use the training part of the dataset to build a decision tree, and then use it to predict the class of a unknown patient, or to prescribe a drug to a new patient. You will use GridSearchCV and Stacking Classifiers to find the best results.

```
[21]: df = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-ML0101EN-SkillsNetwork/labs/Module%203/data/drug200.csv", delimiter=",")
df.head()
```

```
[21]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

Let's create the X and y for our dataset:

```
[22]: X = df[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K']].values
X[0:5]
```

```
[22]: array([[23, 'F', 'HIGH', 'HIGH', 25.355],
        [47, 'M', 'LOW', 'HIGH', 13.093],
        [47, 'M', 'LOW', 'HIGH', 10.114],
        [28, 'F', 'NORMAL', 'HIGH', 7.798],
        [61, 'F', 'LOW', 'HIGH', 18.043]], dtype=object)
```

```
[23]: y = df["Drug"]
y[0:5]
```

```
[23]: 0    drugY
      1    drugC
      2    drugC
      3    drugX
      4    drugY
      Name: Drug, dtype: object
```

Now lets use a LabelEncoder to turn categorical features into numerical:

```
[24]: from sklearn import preprocessing
      le_sex = preprocessing.LabelEncoder()
      le_sex.fit(['F','M'])
      X[:,1] = le_sex.transform(X[:,1])

      le_BP = preprocessing.LabelEncoder()
      le_BP.fit([ 'LOW', 'NORMAL', 'HIGH'])
      X[:,2] = le_BP.transform(X[:,2])

      le_Chol = preprocessing.LabelEncoder()
      le_Chol.fit([ 'NORMAL', 'HIGH'])
      X[:,3] = le_Chol.transform(X[:,3])

      X[0:5]
```

```
[24]: array([[23, 0, 0, 0, 25.355],
             [47, 1, 1, 0, 13.093],
             [47, 1, 1, 0, 10.114],
             [28, 0, 2, 0, 7.798],
             [61, 0, 1, 0, 18.043]], dtype=object)
```

```
[25]: scaler = preprocessing.StandardScaler().fit(X)
      X= scaler.transform(X)
```

Split the data into training and testing data with a 80/20 split

```
[26]: X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.3,
      ↪random_state=4)
      print ('Train set:', X_train.shape,  y_train.shape)
      print ('Test set:', X_test.shape,  y_test.shape)
```

```
Train set: (140, 5) (140,)
Test set: (60, 5) (60,)
```

We have our dictionary of estimators, the individual model objects or base learners.

```
[27]: estimators =_
      ↪[('SVM',SVC(random_state=42)),('knn',KNeighborsClassifier()),('dt',DecisionTreeClassifier())]
      estimators
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_1275/1530583448.py in <module>
----> 1 estimators =_
      ↪[('SVM',SVC(random_state=42)),('knn',KNeighborsClassifier()),('dt',DecisionTreeClassifier())]
      2 estimators

NameError: name 'SVC' is not defined
```

We create a Stacking Classifier:

```
[ ]: clf = StackingClassifier( estimators=estimators, final_estimator=_
      ↪LogisticRegression())
      clf
```

In order to alter the base models in the dictionary of hyperparameter values, we add the key value of each model followed by the parameter of the model we would like to vary.

```
[ ]: param_grid = {'dt__max_depth': [n for n in range(10)], 'dt__random_state':
      ↪[0], 'SVM__C': [0.01, 0.1, 1], 'SVM__kernel': ['linear', 'poly', _
      ↪'rbf'], 'knn__n_neighbors': [1, 4, 8, 9] }
```

We use GridSearchCV to search over specified parameter values of the model.

```
[ ]: search = GridSearchCV(estimator=clf, param_grid=param_grid, scoring='accuracy')
      search.fit(X_train, y_train)
```

We can find the accuracy of the best model.

```
[ ]: search.best_score_
```

We can find the best parameter values:

```
[ ]: search.best_params_
```

We can find the accuracy test data:

```
[ ]: get_accuracy(X_train, X_test, y_train, y_test, search)
```

```
[ ]: # We use sklearn version 0.20.1 for all other labs, please run this command_
      ↪after finishing the lab
```

```
!mamba install -c conda-forge -qy scikit-learn=0.20.1
```

1.1.3 Thank you for completing this lab!

1.2 Author

Joseph Santarcangelo

1.2.1 Other Contributors

1.3 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-01-01	1.0	Joseph S	Created the initial version
2022-02-09	1.1	Steve Hord	QA pass

##

© IBM Corporation 2021. All rights reserved.