

# 04c\_LAB\_Clustering\_Methods

May 24, 2022

## 1 Machine Learning Foundation

### 1.1 Course 4, Part c: Clustering Methods LAB

## 2 Clustering Methods Exercises

### 2.1 Introduction

This lab uses a dataset on wine quality. The data set contains various chemical properties of wine, such as acidity, sugar, pH, and alcohol. It also contains a quality metric (3-9, with highest being better) and a color (red or white). The name of the file is `Wine_Quality_Data.csv`.

We will be using the chemical properties (i.e. everything but quality and color) to cluster the wine. Though this is unsupervised learning, there are interesting semi-supervised extensions relating clustering results onto color and quality.

[ ]:

```
[3]: def warn(*args, **kwargs):  
      pass  
      import warnings  
      warnings.warn = warn  
  
      import seaborn as sns, pandas as pd, numpy as np
```

```
[6]: import numpy as np, pandas as pd, matplotlib.pyplot as plt, seaborn as sns, os
```

```
[5]: from colorspace import colors, palette  
      sns.set_palette(palette)
```

```
-----  
ModuleNotFoundError                                Traceback (most recent call last)  
/tmp/ipykernel_82/3134087529.py in <module>  
      1 import numpy as np, pandas as pd, matplotlib.pyplot as plt, seaborn as   
↪sns, os  
----> 2 from colorspace import colors, palette  
      3 sns.set_palette(palette)
```

```
ModuleNotFoundError: No module named 'colorspace'
```

## 2.2 Question 1

- Import the data and examine the features.
- Note which are continuous, categorical, and boolean.
- How many entries are there for the two colors and range of qualities?
- Make a histogram plot of the quality for each of the wine colors.

```
[7]: ### BEGIN SOLUTION
# Import the data
data = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.
↳ appdomain.cloud/IBM-ML0187EN-SkillsNetwork/labs/module%202/Wine_Quality_Data.
↳ csv")

data.head(4).T
```

```
[7]:
```

	0	1	2	3
fixed_acidity	7.4	7.8	7.8	11.2
volatile_acidity	0.7	0.88	0.76	0.28
citric_acid	0.0	0.0	0.04	0.56
residual_sugar	1.9	2.6	2.3	1.9
chlorides	0.076	0.098	0.092	0.075
free_sulfur_dioxide	11.0	25.0	15.0	17.0
total_sulfur_dioxide	34.0	67.0	54.0	60.0
density	0.9978	0.9968	0.997	0.998
pH	3.51	3.2	3.26	3.16
sulphates	0.56	0.68	0.65	0.58
alcohol	9.4	9.8	9.8	9.8
quality	5	5	5	6
color	red	red	red	red

```
[8]: data.shape
```

```
[8]: (6497, 13)
```

The data types for each entry. The implementation of K-means in Scikit-learn is designed only to work with continuous data (even though it is sometimes used with categorical or boolean types). Fortunately, all the columns we will be using (everything except quality and color) are continuous.

```
[9]: data.dtypes
```

```
[9]: fixed_acidity      float64
volatile_acidity     float64
citric_acid          float64
residual_sugar       float64
chlorides            float64
```

```

free_sulfur_dioxide    float64
total_sulfur_dioxide   float64
density                float64
pH                    float64
sulphates              float64
alcohol                float64
quality                int64
color                  object
dtype: object

```

The number of entries for each wine color.

```
[10]: data.color.value_counts()
```

```

[10]: white    4898
      red      1599
      Name: color, dtype: int64

```

The distribution of quality values.

```
[11]: data.quality.value_counts().sort_index()
```

```

[11]: 3      30
      4     216
      5    2138
      6    2836
      7    1079
      8     193
      9        5
      Name: quality, dtype: int64

```

Now for the histogram.

```

[12]: # seaborn styles
      sns.set_context('notebook')
      sns.set_style('white')

      # custom colors
      red = sns.color_palette()[2]
      white = sns.color_palette()[4]

      # set bins for histogram
      bin_range = np.array([3, 4, 5, 6, 7, 8, 9])

      # plot histogram of quality counts for red and white wines
      ax = plt.axes()
      for color, plot_color in zip(['red', 'white'], [red, white]):
          q_data = data.loc[data.color==color, 'quality']

```


```

q_data.hist(bins=bin_range,
            alpha=0.5, ax=ax,
            color=plot_color, label=color)

ax.legend()
ax.set(xlabel='Quality', ylabel='Occurrence')

# force tick labels to be in middle of region
ax.set_xlim(3,10)
ax.set_xticks(bin_range+0.5)
ax.set_xticklabels(bin_range);
ax.grid('off')
### END SOLUTION

```



## 2.3 Question 2

- Examine the correlation and skew of the relevant variables—everything except color and quality (without dropping these columns from our data).
- Perform any appropriate feature transformations and/or scaling.
- Examine the pairwise distribution of the variables with pairplots to verify scaling and normalization efforts.

```
[13]: ### BEGIN SOLUTION
float_columns = [x for x in data.columns if x not in ['color', 'quality']]

# The correlation matrix
corr_mat = data[float_columns].corr()

# Strip out the diagonal values for the next step
for x in range(len(float_columns)):
    corr_mat.iloc[x,x] = 0.0

corr_mat
```

```
[13]:
```

	fixed_acidity	volatile_acidity	citric_acid	\
fixed_acidity	0.000000	0.219008	0.324436	
volatile_acidity	0.219008	0.000000	-0.377981	
citric_acid	0.324436	-0.377981	0.000000	
residual_sugar	-0.111981	-0.196011	0.142451	
chlorides	0.298195	0.377124	0.038998	
free_sulfur_dioxide	-0.282735	-0.352557	0.133126	
total_sulfur_dioxide	-0.329054	-0.414476	0.195242	
density	0.458910	0.271296	0.096154	
pH	-0.252700	0.261454	-0.329808	
sulphates	0.299568	0.225984	0.056197	
alcohol	-0.095452	-0.037640	-0.010493	

	residual_sugar	chlorides	free_sulfur_dioxide	\
fixed_acidity	-0.111981	0.298195	-0.282735	
volatile_acidity	-0.196011	0.377124	-0.352557	
citric_acid	0.142451	0.038998	0.133126	
residual_sugar	0.000000	-0.128940	0.402871	
chlorides	-0.128940	0.000000	-0.195045	
free_sulfur_dioxide	0.402871	-0.195045	0.000000	
total_sulfur_dioxide	0.495482	-0.279630	0.720934	
density	0.552517	0.362615	0.025717	
pH	-0.267320	0.044708	-0.145854	
sulphates	-0.185927	0.395593	-0.188457	
alcohol	-0.359415	-0.256916	-0.179838	

	total_sulfur_dioxide	density	pH	sulphates	\
fixed_acidity	-0.329054	0.458910	-0.252700	0.299568	
volatile_acidity	-0.414476	0.271296	0.261454	0.225984	
citric_acid	0.195242	0.096154	-0.329808	0.056197	
residual_sugar	0.495482	0.552517	-0.267320	-0.185927	
chlorides	-0.279630	0.362615	0.044708	0.395593	
free_sulfur_dioxide	0.720934	0.025717	-0.145854	-0.188457	
total_sulfur_dioxide	0.000000	0.032395	-0.238413	-0.275727	
density	0.032395	0.000000	0.011686	0.259478	

pH	-0.238413	0.011686	0.000000	0.192123
sulphates	-0.275727	0.259478	0.192123	0.000000
alcohol	-0.265740	-0.686745	0.121248	-0.003029

	alcohol
fixed_acidity	-0.095452
volatile_acidity	-0.037640
citric_acid	-0.010493
residual_sugar	-0.359415
chlorides	-0.256916
free_sulfur_dioxide	-0.179838
total_sulfur_dioxide	-0.265740
density	-0.686745
pH	0.121248
sulphates	-0.003029
alcohol	0.000000

```
[14]: # Pairwise maximal correlations
corr_mat.abs().idxmax()
```

```
[14]: fixed_acidity          density
volatile_acidity    total_sulfur_dioxide
citric_acid          volatile_acidity
residual_sugar          density
chlorides            sulphates
free_sulfur_dioxide    total_sulfur_dioxide
total_sulfur_dioxide    free_sulfur_dioxide
density              alcohol
pH                  citric_acid
sulphates            chlorides
alcohol              density
dtype: object
```

And an examination of the skew values in anticipation of transformations.

```
[15]: skew_columns = (data[float_columns]
                    .skew()
                    .sort_values(ascending=False))

skew_columns = skew_columns.loc[skew_columns > 0.75]
skew_columns
```

```
[15]: chlorides          5.399828
sulphates           1.797270
fixed_acidity       1.723290
volatile_acidity     1.495097
residual_sugar      1.435404
```

```
free_sulfur_dioxide    1.220066
dtype: float64
```

```
[16]: # Perform log transform on skewed columns
      for col in skew_columns.index.tolist():
          data[col] = np.log1p(data[col])
```

Perform feature scaling.

```
[17]: from sklearn.preprocessing import StandardScaler

      sc = StandardScaler()
      data[float_columns] = sc.fit_transform(data[float_columns])

      data.head(4)
```

```
[17]:   fixed_acidity  volatile_acidity  citric_acid  residual_sugar  chlorides \
0      0.229509      2.135767      -2.192833      -0.815173      0.624554
1      0.550261      3.012817      -2.192833      -0.498175      1.281999
2      0.550261      2.438032      -1.917553      -0.625740      1.104012
3      2.802728      -0.337109      1.661085      -0.815173      0.594352


      free_sulfur_dioxide  total_sulfur_dioxide  density      pH  sulphates \
0      -1.193601      -1.446359  1.034993  1.813090  0.250355
1      -0.013944      -0.862469  0.701486 -0.115073  1.059213
2      -0.754684      -1.092486  0.768188  0.258120  0.862549
3      -0.574982      -0.986324  1.101694 -0.363868  0.389396

      alcohol  quality  color
0 -0.915464      5    red
1 -0.580068      5    red
2 -0.580068      5    red
3 -0.580068      6    red
```

Finally, the pairplot of the transformed and scaled features.

```
[18]: sns.set_context('notebook')
      sns.pairplot(data[float_columns + ['color']],
                    hue='color',
                    hue_order=['white', 'red'],
                    palette={'red':red, 'white':'gray'});

      ### END SOLUTION
```



## 2.4 Question 3

- Fit a K-means clustering model with two clusters.
- Examine the clusters by counting the number of red and white wines in each cluster.

```
[19]: from sklearn.cluster import KMeans
      ### BEGIN SOLUTION
      km = KMeans(n_clusters=2, random_state=42)
      km = km.fit(data[float_columns])

      data['kmeans'] = km.predict(data[float_columns])
```

```
[20]: (data[['color', 'kmeans']]
      .groupby(['kmeans', 'color'])
```



```
.size()
.to_frame()
.rename(columns={0: 'number'}))
### END SOLUTION
```

```
[20]:
```

		number
kmeans	color	
0	red	1576
	white	87
1	red	23
	white	4811

## 2.5 Question 4

- Now fit K-Means models with cluster values ranging from 1 to 20.
- For each model, store the number of clusters and the inertia value.
- Plot cluster number vs inertia. Does there appear to be an ideal cluster number?


```
[21]: ### BEGIN SOLUTION
# Create and fit a range of models
km_list = list()

for clust in range(1,21):
    km = KMeans(n_clusters=clust, random_state=42)
    km = km.fit(data[feature_columns])

    km_list.append(pd.Series({'clusters': clust,
                              'inertia': km.inertia_,
                              'model': km}))
```

```
[22]: plot_data = (pd.concat(km_list, axis=1)
                    .T
                    [['clusters', 'inertia']]
                    .set_index('clusters'))

ax = plot_data.plot(marker='o', ls='--')
ax.set_xticks(range(0,21,2))
ax.set_xlim(0,21)
ax.set(xlabel='Cluster', ylabel='Inertia');
### END SOLUTION
```



## 2.6 Question 5

- Fit an agglomerative clustering model with two clusters.
- Compare the results to those obtained by K-means with regards to wine color by reporting the number of red and white observations in each cluster for both K-means and agglomerative clustering.
- Visualize the dendrogram produced by agglomerative clustering. *Hint:* SciPy has a module called `cluster.hierarchy` that contains the `linkage` and `dendrogram` functions required to create the linkage map and plot the resulting dendrogram.

```
[23]: from sklearn.cluster import AgglomerativeClustering
      ### BEGIN SOLUTION
      ag = AgglomerativeClustering(n_clusters=2, linkage='ward',
      ↪compute_full_tree=True)
      ag = ag.fit(data[float_columns])
      data['agglom'] = ag.fit_predict(data[float_columns])
```

Note that cluster assignment is arbitrary, the respective primary cluster numbers for red and white may not be identical to the ones below and also may not be the same for both K-means and agglomerative clustering.

```
[24]: # First, for Agglomerative Clustering:
      (data[['color', 'agglom', 'kmeans']]
       .groupby(['color', 'agglom'])
       .size())
```

```
.to_frame()
.rename(columns={0: 'number'}))
```

```
[24]:          number
color agglom
red   0         31
      1        1568
white 0        4755
      1         143
```

```
[25]: # Comparing with KMeans results:
(data[['color', 'agglom', 'kmeans']]
 .groupby(['color', 'kmeans'])
 .size()
 .to_frame()
 .rename(columns={0: 'number'}))
```

```
[25]:          number
color kmeans
red   0        1576
      1         23
white 0         87
      1       4811
```

```
[26]: # Comparing results:
(data[['color', 'agglom', 'kmeans']]
 .groupby(['color', 'agglom', 'kmeans'])
 .size()
 .to_frame()
 .rename(columns={0: 'number'}))
```

```
[26]:          number
color agglom kmeans
red   0         0         13
      1         1         18
      1         0       1563
      1         1         5
white 0         0         38
      1         1      4717
      1         0         49
      1         1         94
```

Though the cluster numbers are not identical, the clusters are very consistent within a single wine variety (red or white).

And here is a plot of the dendrogram created from agglomerative clustering.

```
[27]: # First, we import the cluster hierarchy module from SciPy (described above) to
      ↪ obtain the linkage and dendrogram functions.
      from scipy.cluster import hierarchy

      Z = hierarchy.linkage(ag.children_, method='ward')

      fig, ax = plt.subplots(figsize=(15,5))

      # Some color setup
      red = colors[2]
      blue = colors[0]


      hierarchy.set_link_color_palette([red, 'gray'])

      den = hierarchy.dendrogram(Z, orientation='top',
                                p=30, truncate_mode='lastp',
                                show_leaf_counts=True, ax=ax,
                                above_threshold_color=blue)

      ### END SOLUTION
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_82/2497210886.py in <module>
      7
      8 # Some color setup
----> 9 red = colors[2]
     10 blue = colors[0]
     11

NameError: name 'colors' is not defined
```



## 2.7 Question 6

In this question, we are going to explore clustering as a form of feature engineering.

- Create a **binary** target variable `y`, denoting if the quality is greater than 7 or not.
- Create a variable called `X_with_kmeans` from `data`, by dropping the columns “quality”, “color” and “agglom” from the dataset. Create `X_without_kmeans` from that by dropping “kmeans”.
- For both datasets, using `StratifiedShuffleSplit` with 10 splits, fit 10 Random Forest Classifiers and find the mean of the ROC-AUC scores from these 10 classifiers.
- Compare the average roc-auc scores for both models, the one using the KMeans cluster as a feature and the one that doesn’t use it.

```
[28]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score
from sklearn.model_selection import StratifiedShuffleSplit

### BEGIN SOLUTION
y = (data['quality'] > 7).astype(int)
X_with_kmeans = data.drop(['agglom', 'color', 'quality'], axis=1)
X_without_kmeans = X_with_kmeans.drop('kmeans', axis=1)
sss = StratifiedShuffleSplit(n_splits=10, random_state=6532)

def get_avg_roc_10splits(estimator, X, y):
    roc_auc_list = []
    for train_index, test_index in sss.split(X, y):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]
        estimator.fit(X_train, y_train)
        y_predicted = estimator.predict(X_test)
        y_scored = estimator.predict_proba(X_test)[: , 1]
        roc_auc_list.append(roc_auc_score(y_test, y_scored))
    return np.mean(roc_auc_list)
# return classification_report(y_test, y_predicted)

estimator = RandomForestClassifier()
roc_with_kmeans = get_avg_roc_10splits(estimator, X_with_kmeans, y)
roc_without_kmeans = get_avg_roc_10splits(estimator, X_without_kmeans, y)
print("Without kmeans cluster as input to Random Forest, roc-auc is \"{0}\"".
      ↪format(roc_without_kmeans))
print("Using kmeans cluster as input to Random Forest, roc-auc is \"{0}\"".
      ↪format(roc_with_kmeans))
```

Without kmeans cluster as input to Random Forest, roc-auc is  
"0.8139007936507937"

Using kmeans cluster as input to Random Forest, roc-auc is "0.8475119047619047"

Let’s now explore if the number of clusters have an effect in this improvement.

- Create the basis training set from `data` by restricting to `float_columns`.

- For  $n = 1, \dots, 20$ , fit a KMeans algorithm with  $n$  clusters. **One-hot encode** it and add it to the **basis** training set. Don't add it to the previous iteration.
- Fit 10 **Logistic Regression** models and compute the average roc-auc-score.
- Plot the average roc-auc scores.

```
[29]: from sklearn.linear_model import LogisticRegression

X_basis = data[float_columns]
sss = StratifiedShuffleSplit(n_splits=10, random_state=6532)

def create_kmeans_columns(n):
    km = KMeans(n_clusters=n)
    km.fit(X_basis)
    km_col = pd.Series(km.predict(X_basis))
    km_cols = pd.get_dummies(km_col, prefix='kmeans_cluster')
    return pd.concat([X_basis, km_cols], axis=1)

estimator = LogisticRegression()
ns = range(1, 21)
roc_auc_list = [get_avg_roc_10splits(estimator, create_kmeans_columns(n), y)
                 for n in ns]

ax = plt.axes()
ax.plot(ns, roc_auc_list)
ax.set(
    xticklabels= ns,
    xlabel='Number of clusters as features',
    ylabel='Average ROC-AUC over 10 iterations',
    title='KMeans + LogisticRegression'
)
ax.grid(True)
### END SOLUTION
```



---

2.7.1 Machine Learning Foundation (C) 2020 IBM Corporation